



吉林大学

学习汇报

汇 报 人：高金彤

2022.12.8

The Majority Can Help the Minority: Context-rich Minority Oversampling for Long-tailed Classification

Seulki Park^{1*} Youngkyu Hong² Byeongho Heo² Sangdoo Yun² Jin Young Choi¹

¹ASRI, ECE., Seoul National University ²NAVER AI Lab

`seulki.park@snu.ac.kr, yk.hong@kaist.ac.kr, {bh.heo, sangdoo.yun}@navercorp.com, jychoi@snu.ac.kr`

长尾分类的根本问题：类数据不平衡数据集由于缺乏少数类的数据，分类器的泛化性能变差

本篇文章的问题：少样本背景单一

关键思路：少样本多样化（多数类帮助少数类——提供背景实现多样性）

The Majority Can Help the Minority

（少数群体过采样方法，利用多样本的丰富图像作为背景来增强少样本的多样化）

提出：利用Cutmix混合手段生成新的以少样本为中心的图像，具有多样本的背景。

主要创新点：少样本加权分布 Q 、权重值、即插即用、计算成本低



Random oversampling



Context-rich oversampling

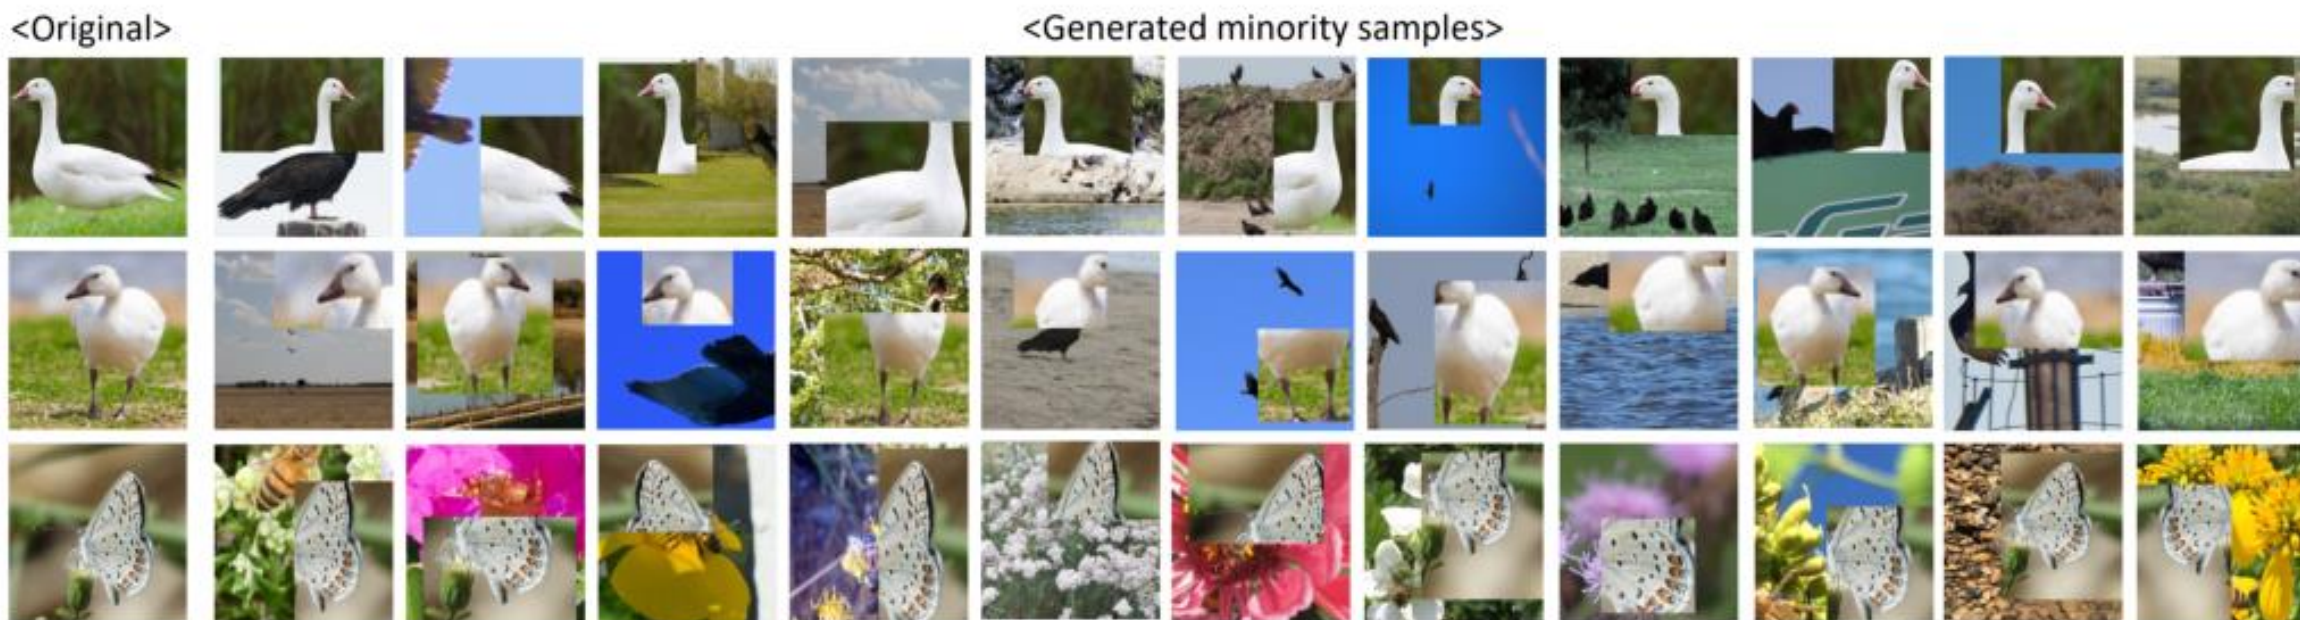


Figure 2. **A display of the minority images generated by CMO** (minority classes: the snow goose and the Acmon blue (butterfly)). We randomly choose generated images for each original image. Our method is able to generate context-rich minority samples that have diverse contexts. For example, while the original ‘snow goose’ class contains only images of a ‘snow goose’ on grass, the generated images have various contexts such as the sky, the sea, the sand, and a flock of crows. These generated images enable the model to learn a robust representation of minority classes.

给定一个来自少数类的原始图像，该对象被裁剪成各种大小，并粘贴到来自多数类的各种图像上。然后，我们可以创建具有更多样化背景图像(例如，带有天空、道路、屋顶、乌鸦等的“雪鹅”图像)。由于这是多数类和少数类样本的插值，因此在决策边界周围生成多样化的数据，从而提高了少数类的泛化性能。

Re-weighting methods: 为训练样本分配不同的权重（概率）





inverse class frequency、CE、LDAM、BS、IB、LADE (loss)

Re-sampling methods: 修改训练分布，以降低不平衡水平（图像数量）

1、输入空间：欠采样、过采样（ROS）、合成少样本过采样（SMOTE）、生成对抗少样本过采样（GAMO）

2、特征空间：深度过采样(DOS)、特征空间增强(FSA)和元语义增强(MetaSAug)

Other long-tailed methods: 延迟重加权(DRW)，分类器再训练(cRT)，可学习权重缩放(LWS)和Mixup移位标签感知平滑模型

	ResNet-50	Mixup [48]	Cutout [3]	CutMix
Image				
Label	Dog 1.0	Dog 0.5 Cat 0.5	Dog 1.0	Dog 0.6 Cat 0.4

知乎 @kaihoo

Algorithm 1 Context-rich Minority Oversampling (CMO)

Require: Dataset $\mathcal{D}_{i=1}^N$, model parameters θ , P , Q , any loss function $L(\cdot)$.

- 1: Randomly initialize θ .
 - 2: Sample weighted dataset $\tilde{\mathcal{D}}_{i=1}^N \sim Q$.
 - 3: **for** epoch = 1, ..., T **do**
 - 4: **for** batch $i = 1, \dots, B$ **do**
 - 5: Draw a mini-batch (x_i^b, y_i^b) from $\mathcal{D}_{i=1}^N$
 - 6: Draw a mini-batch (x_i^f, y_i^f) from $\tilde{\mathcal{D}}_{i=1}^N$
 - 7: $\lambda \sim \text{Beta}(\alpha, \alpha)$
 - 8: $\tilde{x}_i = \mathbf{M} \odot x_i^b + (\mathbf{1} - \mathbf{M}) \odot x_i^f$
 - 9: $\tilde{y}_i = \lambda y_i^b + (1 - \lambda) y_i^f$
 - 10: $\theta \leftarrow \theta - \eta \nabla L((\tilde{x}_i, \tilde{y}_i); \theta)$
 - 11: **end for**
 - 12: **end for**
-

Algorithm 2 PyTorch-style pseudo-code for CMO

```
# original_loader: data loader from original data distribution
# weighted_loader: data loader from minor-class-weighted distribution
# model: any backbone network such as ResNet or multi-branch networks (RIDE)
# loss: any loss such as CE, LDAM, balanced softmax, RIDE loss

for epoch in Epochs:
    # load a batch for background images from original data dist.
    for x_b, y_b in original_loader:
        # load a batch for foreground from minor-class-weighted dist.
        x_f, y_f = next(weighted_loader)

        # get coordinate for random binary mask
        lambda = np.random.uniform(0,1)
        cx = np.random.randint(W) # W: width of images
        cy = np.random.randint(H) # H: height of images
        bbx1 = np.clip(cx - int(W * np.sqrt(1. - lambda))/2, 0, W)
        bbx2 = np.clip(cx + int(W * np.sqrt(1. - lambda))/2, 0, W)
        bby1 = np.clip(cy - int(H * np.sqrt(1. - lambda))/2, 0, H)
        bby2 = np.clip(cy + int(H * np.sqrt(1. - lambda))/2, 0, H)

        # get minor-oversampled images
        x_b[:, :, bbx1:bbx2, bby1:bby2] = x_f[:, :, bbx1:bbx2, bby1:bby2]
        lambda = 1 - ((bbx2 - bbx1) * (bby2 - bby1) / (W * H)) # adjust lambda

        # output (x_f is attached to x_b)
        output = model(x_b)

        # loss
        losses = loss(output, y_b) * lambda + loss(output, y_f) * (1. - lambda)

        # optimization step
        losses.backward()
        optimizer.step()
```

Let n_k be the number of samples in the k -th class, then for the C classes, the total number of samples is $N = \sum_{k=1}^C n_k$. Then, the generalized sampling probability for the k -th class can be defined by

$$q(r, k) = \frac{1/n_k^r}{\sum_{k'=1}^C 1/n_{k'}^r}, \quad (2)$$

$r = 1$ the inverse class frequency

$r = 1/2$ the smoothed inverse class frequency

```
def get_weighted_sampler(self):
    cls_num_list = self.get_cls_num_list()
    cls_weight = 1.0 / (np.array(cls_num_list) ** self.weighted_alpha)
    cls_weight = cls_weight / np.sum(cls_weight) * len(cls_num_list)
    samples_weight = np.array([cls_weight[t] for t in self.targets])
    samples_weight = torch.from_numpy(samples_weight)
    samples_weight = samples_weight.double()
    print("samples_weight", samples_weight)
    sampler = torch.utils.data.WeightedRandomSampler(samples_weight, len(self.targets), replacement=True)
    return sampler
```


$$\begin{aligned}\tilde{x} &= \mathbf{M} \odot x^b + (\mathbf{1} - \mathbf{M}) \odot x^f \\ \tilde{y} &= \lambda y^b + (1 - \lambda)y^f,\end{aligned}\tag{1}$$

缓解过度置信度问题：公式1生成的标签会惩罚过度自信的输出，类似于标签平滑正则化

Table 1. **Summary of datasets.** The imbalance ratio ρ is defined by $\rho = \max_k \{n_k\} / \min_k \{n_k\}$, where n_k is the number of samples in the k -th class.

Dataset	# of classes	# of training	Imbalance ratio
CIFAR-100-LT	100	50K	$\{10, 50, 100\}$
ImageNet-LT	1,000	115.8K	256
iNaturalist 2018	8,142	437.5K	500

不平衡比率=100（最大类的数量/最小类的数量=500/5=100）：一共有10847张图像

imb_factor=0.01|

[500, 477, 455, 434, 415, 396, 378, 361, 344, 328, 314, 299, 286, 273,

不平衡比率=10（最大类的数量/最小类的数量=500/50=10）：一共有19573张图像

imb_factor=0.1

[500, 488, 477, 466, 455, 445, 434, 424, 415, 405, 396, 387, 378, 369,
361, 352, 344, 336, 328, 321, 314, 306, 299, 292, 286, 279, 273, 266,
260, 254, 248, 243, 237, 232, 226, 221, 216, 211, 206, 201, 197, 192,
188, 183, 179, 175, 171, 167, 163, 159, 156, 152, 149, 145, 142, 139,
135, 132, 129, 126, 123, 121, 118, 115, 112, 110, 107, 105, 102, 100, 98,
95, 93, 91, 89, 87, 85, 83, 81, 79, 77, 75, 74, 72, 70, 69, 67, 66, 64,
63, 61, 60, 58, 57, 56, 54, 53, 52, 51, 50]
287, 276, 265, 255, 245, 235, 226, 218, 209, 201, 193, 186, 178, 172,
165, 158, 152, 146, 141, 135, 130, 125, 120, 115, 111, 107, 102, 98, 95,
91, 87, 84, 81, 78, 75, 72, 69, 66, 64, 61, 59, 56, 54, 52, 50, 48, 46,
44, 43, 41, 39, 38, 36, 35, 34, 32, 31, 30, 29, 27, 26, 25, 24, 23, 22,
22, 21, 20, 19, 18, 18, 17, 16, 16, 15, 14, 14, 13, 13, 12, 12, 11, 11,
10, 10, 10]

设备:	单个NVIDIA GTX 1080Ti GPU
图片大小	32*32
框架与优化器:	Pytorch、SGD(momentum:0.9 weight decay: $2*10^{-4}$)
Backbone	Resnet-32
数据增强:	每边填充4个像素，并将水平翻转或随机裁剪应用到 32×32 大小
学习率:	初始值为0.1，在第160和180个时期*0.01 在前五个epoch使用了学习率的线性预热linear warm-up of the learning rate
Epoch, Batchsize	200, 128
评价指标:	Top1 Accuracy(overall, many, medium, few)
备注:	为了在原始输入空间中微调模型，我们在最后三个阶段关闭了CMO。 我发现在前3个epoch也关闭了CMO

设备:	4个NVIDIA GTX 1080Ti GPU
图片大小	224*224
框架与优化器:	Pytorch、SGD(momentum:0.9)
Backbone	Resnet-50
数据增强:	水平翻转、颜色抖动
学习率:	初始值为0.1, 在第60和80个时期*0.1
Epoch, Batchsize	100, 256
评价指标:	Top1 Accuracy(overall, many, medium, few)
备注:	为了在原始输入空间中微调模型, 我们在最后三个阶段关闭了CMO。 我发现在前3个epoch也关闭了CMO

设备:	8个Tesla V100 GPU
框架与优化器:	Pytorch、SGD(momentum:0.9)
Backbone	Resnet-50、101、152和Wide ResNet-50
数据增强:	水平翻转、颜色抖动
学习率:	初始值为0.1，在第75和160个时期*0.1
Epoch, Batchsize	200, 512
评价指标:	Top1 Accuracy(overall, many, medium, few)
备注:	实验在NAVER智能机器学习（NSML）平台上实施和评估 为了在原始输入空间中微调模型，我们在最后三个阶段关闭了CMO。 我发现在前3个epoch也关闭了CMO

Table 2. **State-of-the-art comparison on CIFAR-100-LT dataset.** Classification accuracy (%) for ResNet-32 architecture on CIFAR-100-LT with different imbalance ratios. * and † are from the original paper.

Table 13. **Comparison against baselines on CIFAR-100-LT** Results with classification accuracy (%) of ResNet-32. The best results are marked in bold.

LT (Im-Net-32.

Imbalance ratio	50				10			
	Vanilla	+ROS [47]	+Remix [7]	+CMO	Vanilla	+ROS [47]	+Remix [7]	+CMO
CE	44.0 (+0.0)	39.7 (-4.3)	45.0 (+1.0)	48.3 (+4.3)	56.4 (+0.0)	55.6 (-0.8)	58.7 (+2.3)	59.5 (+3.1)
CE-DRW [5]	45.6 (+0.0)	41.3 (-4.3)	49.5 (+3.9)	50.9 (+5.3)	57.9 (+0.0)	56.4 (-1.5)	59.2 (+1.3)	61.7 (+3.8)
LDAM-DRW [5]	47.9 (+0.0)	38.3 (-9.6)	48.8 (+0.9)	51.7 (+3.8)	57.3 (+0.0)	53.9 (-3.4)	55.9 (-1.4)	58.4 (+1.1)
RIDE [49]	51.4 (+0.0)	31.3 (-20.1)	47.9 (-3.5)	53.0 (+1.6)	59.8 (+0.0)	49.4 (-10.4)	59.5 (-0.3)	60.2 (+0.4)

MiSLAS [56]*

CE + CMO

CE-DRW + CMO

LDAM-DRW + CMO

BS + CMO

RIDE (3 experts) + CMO

47.0

52.3

63.2

43.9

48.3

59.5

47.0

50.9

61.7

47.2

51.7

58.4

46.6

51.4

62.3

50.0

53.0

60.2

Table 4. **State-of-the-art comparison on ImageNet-LT.** Classification accuracy (%) of ResNet-50 with state-of-the-art methods trained for 90 or 100 epochs. “*” and “†” denote the results are from the original papers, and [25], respectively. The best results are marked in bold.

	All	Many	Med	Few
Cross Entropy (CE) [†]	41.6	64.0	33.8	5.8
Decouple-cRT [25] [†]	47.3	58.8	44.0	26.1
Decouple-LWS [25] [†]	47.7	57.1	45.2	29.3
Remix [7]	48.6	60.4	46.9	30.7
LDAM-DRW [5]	49.8	60.4	46.9	30.7
CE-DRW	50.1	61.7	47.3	28.8
Balanced Softmax (BS) [40]	51.0	60.9	48.8	32.1
Causal Norm [46]*	51.8	62.7	48.8	31.6
RIDE (3 experts) [49]*	54.9	66.2	51.7	34.9
RIDE (4 experts) [49]*	55.4	66.2	52.3	36.5
CE + CMO	49.1	67.0	42.3	20.5
CE-DRW + CMO	51.4	60.8	48.6	35.5
LDAM-DRW + CMO	51.1	62.0	47.4	30.8
BS + CMO	52.3	62.0	49.1	36.7
RIDE (3 experts) + CMO	56.2	66.4	53.9	35.6

Table 5. **Comparison against baselines on ImageNet-LT.** Classification accuracy (%) of ResNet-50.

	Vanilla	+Remix [7]	+CMO
CE	41.6 (+0.0)	41.7 (+0.1)	49.1 (+7.5)
CE-DRW [5]	50.1 (+0.0)	48.6 (-1.5)	51.4 (+1.3)
Balanced Softmax [40]	51.0 (+0.0)	49.2 (-1.8)	52.3 (+1.3)

Table 7. **State-of-the-art comparison on iNaturalist2018.** Classification accuracy (%) of ResNet-50 on iNaturalist2018. “*” and “†” indicate the results from the original paper and [57], respectively. RIDE [49] was trained for 100 epochs.

	All	Many	Med	Few
Cross Entropy (CE)	61.0	73.9	63.5	55.5
IB Loss [38]*	65.4	-	-	-
FSA [8]*	65.9	-	-	-
LDAM-DRW [5]†	66.1	-	-	-
Decouple-cRT [25]*	68.2	73.2	68.8	66.1
Decouple-LWS [25]*	69.5	71.0	69.8	68.8
BBN [57]*	69.6	-	-	-
Balanced Softmax [40]	70.0	70.0	70.2	69.9
LADE [20]*	70.0	-	-	-
Remix [7]*	70.5	-	-	-
MiSLAS [56]*	71.6	73.2	72.4	70.4
RIDE (3 experts) [49]*	72.2	70.2	72.2	72.7
RIDE (4 experts) [49]*	72.6	70.9	72.4	73.1
CE + CMO	68.9	76.9	69.3	66.6
CE-DRW + CMO	70.9	68.2	70.2	72.2
LDAM-DRW + CMO	69.1	75.3	69.5	67.3
BS + CMO	70.9	68.8	70.0	72.3
CE-DRW + CMO + LAS [56]	71.8	69.6	72.1	71.9
RIDE (3 experts) + CMO	72.8	68.7	72.6	73.1

Table 8. **Results on large architectures.** Classification accuracy (%) of large backbone networks on iNaturalist 2018. The results are copied from [8].

Method	ResNet-50	Wide ResNet-50	ResNet-101	ResNet-152
CE	61.0	-	65.2	66.2
FSA [8]	65.9	-	68.4	69.1
CMO	70.9	71.9	72.4	72.6

Imbalance ratio	100	50	10
BS*	50.8	54.2	63.0
PaCo [11]*	52.0	56.0	64.2
BS + CMO	51.7	56.7	65.3

Table 14. **Classification Accuracy on CIFAR-100-LT with different imbalance ratios.** We train ResNet-32 with AutoAugment [9] in 400 epochs. * is from [11] The best results are marked in bold.

Table 6. **Results on longer training epochs with RandAugment [10].** Classification accuracy (%) of ResNet-50 on ImageNet-LT. “*” denotes the results from [11].

	All	Many	Med	Few
BS*	55.0	66.7	52.9	33.0
PaCo [11]*	57.0	65.0	55.7	38.2
BS + CMO	58.0	67.0	55.0	44.2

	All	Many	Med	Few
BS*	71.8	72.3	72.6	71.7
PaCo [11]*	73.2	70.3	73.2	73.6
BS + CMO	74.0	71.9	74.2	74.2

Table 15. **Classification Accuracy on iNaturalist2018.** We train ResNet-50 for 400 epochs with RandAugment [10]. “*” indicates the results are from [11]. The best results are marked in bold.




Figure 3. **Impact of α** on CIFAR-100-LT with an imbalance ratio of 100.

Is the distribution for augmenting images important? (比Cutmix好的原因)

How to choose the appropriate probability distribution Q ? (Q 的 r 选择)

Why should we oversample only for the foreground samples? (只在前景使用的原因)

Comparison with other minority augmentations. (常见数据增强的手段)

Table 9. **Comparison with CutMix** using cross-entropy loss.

	All	Many	Med	Few
<i>CIFAR-100-LT</i>				
CutMix	35.6	71.0	37.9	4.9
CMO	43.9	70.4	42.5	14.4
<i>ImageNet-LT</i>				
CutMix	45.5	68.6	38.1	8.1
CMO	49.1	67.0	42.3	20.5

Table 10. **Impact of different Q sampling distributions.** Results on CIFAR-100-LT (imbalance ratio=100) according to different Q sampling probabilities.

	All	Many	Med	Few
$q(1/2, k)$	42.6	71.6	42.1	9.5
$q(1, k)$	43.9	70.4	42.5	14.4
$q(2, k)$	40.1	67.2	36.7	12.3
$E(k)$ [12]	39.5	70.4	38.0	4.7

Why should we oversample only for the foreground samples?



吉林大学

Table 11. **Ablation study.** Results from variants of CMO with ResNet-32 on imbalanced CIFAR-100; imbalance ratio of 100.

	All	Many	Med	Few
Cross Entropy (CE)	38.6	65.3	37.6	8.7
CE + CMO _{minor}	37.9	58.3	40.4	11.2
CE + CMO _{back}	40.1	64.7	40.2	11.3
CE + CMO	43.9	70.4	42.5	14.4
LDAM [5]	41.7	61.4	42.2	18.0
LDAM + CMO _{minor}	31.7	50.2	33.2	8.4
LDAM + CMO _{back}	44.2	59.2	46.6	24.0
LDAM + CMO	47.2	61.5	48.6	28.8

这是因为，使用CutMix方法，前景图像中的对象很有可能与背景图像重叠。因此，我们可以预期背景图像中关于少样本的信息会丢失，从而导致有限的性能提升。

Table 12. **Data augmentation methods.** Comparisons between augmentation methods for generating new minority samples on CIFAR-100-LT with an imbalance ratio of 100.

	All	Many	Med	Few
CMO w/ Gaussian Blur	31.1	54.7	28.8	6.2
CMO w/ Color Jitter	34.7	58.9	34.4	6.8
CMO w/ Mixup	38.0	54.8	40.2	15.9
CMO w/ CutMix	43.9	70.4	42.5	14.4

Q1

少样本和多样本的效果难以实现同时上升。

Q2

硬件问题


I1

分组：500抽5张+少5张->一组

500中初概率相同，被选中该图像的概率下降，未选中概率上升

I2

通过损失函数（多损失、中损失和少损失）+GAN





吉林大学

— THANKS —
恳请老师批评指正

汇报人：高金彤

2022.12.8